

Multimodal Risk Fusion System for Real-Time Driver Safety

MRS PARVATHI K, Dr. ARADHANA D, Dr. JAGADISH R M,
N HARIKA, MD BASHEER KHAN, MALLIKA K, N NIKHIL

Dept. of CSE (Data Science)
Ballari Institute of Technology and Management
Ballari, India

parvathik035@gmail.com, aradhana@bitm.edu.in, rm.jagadish@gmail.com,
nharika9019@gmail.com, basheerkhaleelkhan@gmail.com, mallikakondapur@gmail.com, nikhilos13@gmail.com

Abstract—Road accidents remain a major global concern, with driver impairment and hazardous road conditions being among the leading causes. Conventional driver assistance systems typically operate in isolation, focusing on environmental perception or driver behavior, failing to capture compounded risk scenarios that arise when both factors coexist.

This paper proposes a real-time dual camera multimodal driver safety system that simultaneously monitors road hazards and driver state using two independent vision pipelines. The system employs a custom-trained YOLOv11 model to detect road hazards such as potholes, debris, and pedestrians, while MediaPipe FaceMesh is utilized to extract facial landmarks to detect driver drowsiness through geometric analysis. Additionally, DeepFace, a pre-trained facial analysis framework, is used to identify driver emotional states including anger, fear, and stress.

A novel risk fusion engine is introduced to integrate heterogeneous input from both modules into a unified risk score. The fusion mechanism incorporates a compounding formulation to model the amplified danger when environmental and driver-related risks occur simultaneously. Based on the computed risk level, the system triggers a multi-layer response mechanism including hardware alerts using Arduino, vibration feedback, and AI-generated voice assistance to ensure timely intervention.

The proposed system offers a scalable, cost-effective and real-time solution to improve road safety by combining perception, analysis, and response within a unified framework. This work addresses critical limitations of existing single-domain systems and provides a foundation for next-generation intelligent driver assistance systems.

Index Terms—Driver safety, Multimodal risk fusion, YOLOv11, DeepFace, MediaPipe, Driver monitoring, Real-time system

I. INTRODUCTION

Road accidents continue to pose a significant threat to public safety, resulting in substantial loss of life and property worldwide. A large proportion of these accidents are attributed to driver-related factors such as fatigue, distraction, and emotional instability, as well as environmental factors that include poor road conditions, obstacles, and pedestrian movement. These factors often occur simultaneously, leading to compound risk scenarios that are difficult to detect using conventional driver assistance systems.

Traditional approaches primarily focus on single-domain monitoring, either analyzing driver behavior or detecting external hazards. Although driver monitoring systems can identify

drowsiness through facial signals such as eye closure and blink rate, object detection models have demonstrated effectiveness in identifying road hazards such as potholes and debris. However, the lack of integration between these two domains limits the ability of such systems to accurately assess real-world risk conditions.

Recent advances in artificial intelligence and computer vision have enabled real-time processing of multiple data streams. Despite this progress, effectively combining heterogeneous inputs into a unified decision-making framework remains a significant challenge. Existing systems fail to capture the compound effect of simultaneous driver impairment and environmental hazards, resulting in incomplete risk evaluation.

To address these limitations, this work proposes a Dual Camera Multimodal Risk Fusion System for real-time driver safety monitoring. The framework utilizes two independent camera streams: a road-facing camera for hazard detection using a YOLOv11-based model, and a driver-facing camera for analyzing driver state using MediaPipe FaceMesh and DeepFace.

The primary contribution of this work is the development of a risk fusion engine that integrates road and driver risks into a unified score. Unlike traditional additive approaches, the proposed model incorporates a compounding mechanism to capture amplified risk when multiple factors coexist. Additionally, the system includes a multi-layer alert mechanism combining hardware-based feedback and AI-driven voice assistance to enable timely intervention.

This framework provides a scalable and real-time solution to improve road safety and contributes to the development of next-generation intelligent driver assistance systems.

II. LITERATURE SURVEY

Significant research has been conducted in the domains of driver monitoring and road hazard detection, using advances in computer vision and deep learning techniques. However, most existing approaches focus on isolated aspects of road safety, limiting their effectiveness in real-world scenarios.

Recent studies have explored the use of object detection models to identify road hazards. Khare et al. [1] proposed a YOLO-based framework to detect road anomalies such as

potholes and obstacles in real time. Their work demonstrated high detection accuracy and low inference time, making it suitable for real-time applications. However, the system is limited to environmental perception and does not incorporate driver behavior analysis, thus lacking a complete safety assessment.

The detection of driver drowsiness has also been extensively studied using facial landmark-based approaches. Das et al. [2] developed a system that combines convolutional neural networks and recurrent architectures to detect fatigue of the driver using temporal movements of patterns in the eyes. Although the model achieved high accuracy under controlled conditions, it requires significant computational resources and is sensitive to environmental variations such as lighting and occlusions.

Emotion recognition in drivers has gained attention as an important factor influencing driving safety. Espino-Salinas et al. [3] utilized deep learning techniques to classify driver emotions using facial expressions, achieving notable accuracy in detecting negative emotional states. Similarly, Venkatesan et al. [4] demonstrated the effectiveness of DeepFace for real-time emotion detection. However, these systems focus solely on emotional analysis and do not integrate other risk factors such as road conditions or physical fatigue.

Several review studies have highlighted the limitations of existing driver monitoring systems. Fonseca and Ferreira [5] conducted a comprehensive survey of deep learning models for the detection of drowsiness and concluded that while deep learning approaches outperform traditional methods, most systems rely on single indicators and lack real-time integration with external environmental factors.

Although these studies demonstrate considerable progress in individual domains, a common limitation among them is the lack of multimodal integration. Most systems operate independently, focusing on driver behavior or environmental perception, and fail to account for the compounded risk that arises when both factors occur simultaneously.

Furthermore, many existing approaches lack real-time response mechanisms, such as multi-layer alert systems or adaptive feedback, which are essential for practical deployment in real-world driving environments.

To address these challenges, this work introduces a multimodal framework that integrates road hazard detection and driver monitoring into a unified architecture. By combining heterogeneous input using a novel risk fusion model, the system provides a more comprehensive and accurate assessment of driving risk. In addition, the inclusion of a multi-layer alert mechanism improves the system's ability to respond effectively to critical situations.

III. METHODOLOGY

A. System Overview

The proposed system follows a multimodal perception and decision pipeline that integrates two independent vision streams to monitor environmental and driver-related risks in real time. The system consists of a road camera for detecting external hazards and a driver-facing camera for analyzing

driver behavior. Each input stream is processed through dedicated modules to extract meaningful features, which are subsequently combined using a risk fusion engine. The final output is a unified risk score that drives a multi-layer alert mechanism.

B. Vision Input Module

1) *Road Camera Pipeline*: The road-facing camera captures continuous video frames that represent the driving environment. Each frame is processed using OpenCV for preprocessing operations such as resizing, normalization, and color space conversion.

$$Frame_{RGB} = Convert(Frame_{BGR}) \quad (1)$$

To ensure real-time performance, frames are resized to a fixed resolution and normalized to reduce computational overhead. These frames are then forwarded to the road hazard detection module.

2) *Driver Camera Pipeline*: The driver-facing camera captures real-time facial data of the driver. The captured frames are processed to detect and localize the facial region.

$$Face_{region} = Crop(Frame, BoundingBox) \quad (2)$$

This step isolates the relevant region of interest, reduces noise, and improves the efficiency of subsequent processing stages.

C. Road Hazard Detection Module (YOLOv11)

The road hazard detection module utilizes a custom-trained YOLOv11 model to identify critical objects such as potholes, road debris, and pedestrians. For each input frame, the model outputs bounding boxes, confidence scores, and class labels. Each detected object is assigned a severity score based on its potential risk.

The road risk score is computed as

$$R_r = \sum_{i=1}^n w_i S_i \quad (3)$$

where S_i represents the severity of the detected object and w_i represents the weight assigned to each class.

This formulation enables the system to prioritize more dangerous objects, such as pedestrians, over minor obstacles.

D. Driver State Monitoring Module

1) *Facial Landmark Detection*: The system employs MediaPipe FaceMesh to extract 468 facial landmarks, which provide detailed geometric information about facial features.

$$L = \{(x_i, y_i)\}, \quad i = 1, 2, \dots, 468 \quad (4)$$

These landmarks serve as the basis for extracting behavioral and fatigue-related features.

2) *Eye Aspect Ratio (EAR) for Drowsiness Detection:* The Eye Aspect Ratio (EAR) is used to estimate eye openness based on selected eye landmarks.

$$EAR = \frac{||P_2 - P_6|| + ||P_3 - P_5||}{2 \cdot ||P_1 - P_4||} \quad (5)$$

A decrease in EAR below a predefined threshold indicates eye closure, which is a strong indicator of drowsiness.

3) *Emotion Detection Using DeepFace:* DeepFace is utilized to classify driver emotions in real time. The model detects emotional states such as anger, fear, and stress. These emotional signals are important indicators of risky driving behavior.

4) *Driver Risk Score Computation:* The driver risk score is computed by combining multiple features:

$$R_d = \alpha F_{drowsiness} + \beta F_{emotion} + \gamma F_{behavior} \quad (6)$$

where α, β, γ are the weighting coefficients, and each feature is normalized to ensure consistent scaling.

E. Risk Fusion Engine (Core Contribution)

To ensure consistent scaling, the computed risk values are normalized:

$$\hat{R} = \frac{R - R_{min}}{R_{max} - R_{min}} \quad (7)$$

This normalization enables for uniform comparison between heterogeneous inputs.

The proposed system introduces a multimodal risk fusion mechanism that integrates road risk and driver risk into a unified score:

$$R = R_r + R_d + 1.5(R_r \cdot R_d) \quad (8)$$

This formulation captures both independent contributions of each risk and compounded risk when both occur simultaneously. Unlike traditional additive models, the multiplicative term amplifies the overall risk, ensuring that dangerous situations are prioritized.

F. Risk Classification

The final risk score is categorized into four levels:

$$R < 30 \rightarrow \text{Safe} \quad (9)$$

$$30 \leq R < 60 \rightarrow \text{Medium} \quad (10)$$

$$60 \leq R < 80 \rightarrow \text{High} \quad (11)$$

$$R \geq 80 \rightarrow \text{Critical} \quad (12)$$

G. Alert and Response System

The system incorporates a multi-layer alert mechanism to ensure effective driver intervention. Medium risk triggers visual alerts using LEDs, high risk activates audible alerts using a buzzer, and critical risk results in vibration motor activation and emergency warnings. The alerts are controlled using an Arduino-based hardware interface, enabling real-time physical feedback.

H. AI-Based Voice Assistance

To improve driver awareness, the system integrates an AI-based voice assistant. The current risk context is passed to a language model that generates concise safety instructions. These instructions are converted to speech using a text-to-speech engine, allowing the driver to receive guidance without diverting attention from the road.

I. System Integration

The system follows a modular data flow architecture, where outputs from individual modules are combined in a centralized fusion engine. This design ensures scalability and allows for independent optimization of each component. The complete pipeline operates in real time, ensuring continuous monitoring and timely intervention.

IV. SYSTEM ARCHITECTURE

The overall architecture of the Dual Camera Multimodal Risk Fusion System follows a structured perception–analysis–decision framework designed for real-time driver safety monitoring. The system utilizes a dual-camera setup, where a road-facing camera captures environmental conditions and a driver-facing camera monitors the driver's behavior. This parallel input structure enables the simultaneous analysis of both external and internal risk factors.

A. Architectural Overview

The road-facing stream is processed using a YOLOv11-based model to detect hazards such as potholes, debris, and pedestrians, which are used to compute a road risk score. In parallel, the driver-facing stream is analyzed using MediaPipe for facial landmark extraction, DeepFace for emotion recognition, and YOLOv8 to detect distraction-related objects such as mobile phones and earpods, resulting in a driver risk score. Both risk scores are integrated using a centralized risk fusion engine to generate a unified risk metric. Based on this final score, the system triggers appropriate alert mechanisms, including visual, auditory, and haptic feedback, along with AI-based voice assistance. The modular architecture ensures scalability, robustness, and efficient real-time performance.

B. Data Flow Pipeline

The system follows a structured data flow pipeline in which inputs from both camera streams are processed in parallel. The road-facing camera captures video frames that are preprocessed and passed to the YOLOv11 model for hazard detection, where detected objects are assigned severity levels




Fig. 1. Overall workflow of the proposed multimodal driver safety system

to compute a road risk score. Simultaneously, the driver-facing camera captures facial data, which is processed using MediaPipe, DeepFace, and YOLOv8 to extract features related to drowsiness, emotion, and distraction, generating a driver risk score. These scores are then integrated using the risk fusion engine to produce a unified risk metric, which determines the appropriate alert response.

C. Flow Diagram Representation

The general workflow of the proposed system is illustrated in Fig. 1.

D. Module Interaction

Each component operates as part of an interconnected pipeline, where outputs from one module serve as inputs to subsequent stages. The road hazard detection module generates a road risk score, while the driver monitoring module produces a driver risk score based on behavioral analysis. These outputs are processed by the risk fusion engine to generate a unified risk score. The final output is passed to the alert and response module, which activates hardware alerts and voice assistance to provide real-time feedback to the driver.

E. Real-Time Processing Considerations

Real-time performance is achieved through the parallel processing of both camera streams, enabling simultaneous analysis of road conditions and driver behavior. Efficiency is maintained through frame preprocessing, optimized models, and streamlined data flow between modules. The system operates at approximately 25–30 frames per second, ensuring timely detection and response.

F. System Characteristics

The framework is designed to be scalable, robust, and cost-effective. Its modular architecture allows independent optimization and easy integration of additional components. The use of efficient models and low-cost hardware enables real-time operation, making it suitable for practical deployment in intelligent driver safety applications.

V. DISCUSSION AND ANALYSIS

A. Technical Feasibility

The system is designed using a combination of custom-trained and pre-trained models to ensure both efficiency and practicality in real-world deployment. The road hazard detection module utilizes a YOLOv11 model optimized for real-time object detection with low latency. For driver monitoring, MediaPipe FaceMesh provides accurate facial landmark detection, while DeepFace enables reliable emotion recognition using pre-trained models. Additionally, YOLOv8 is incorporated to detect distracting objects such as mobile phones and earpods. The integration of these technologies ensures a balance between accuracy, speed, and computational efficiency, making the framework suitable for real-time applications.

B. Hardware Feasibility

The implementation relies on low-cost and easily accessible hardware components, making it suitable for academic and real-world deployment. The setup includes a standard laptop with moderate processing capability, along with two cameras to capture road and driver inputs. An Arduino Nano interfaces with components such as LEDs, a buzzer, and a vibration motor to generate alerts. The design avoids complex circuitry by utilizing simple USB-based connections, reducing installation complexity and improving portability. These minimal hardware requirements enable deployment and testing with limited resources.

C. Computational Feasibility

Efficient computational performance is essential for real-time operation. The framework employs optimized models such as YOLOv11, YOLOv8, and MediaPipe to achieve fast inference with moderate resource requirements. Input frames are resized and normalized to reduce computational load while maintaining accuracy. Parallel processing of road and driver-facing streams minimizes latency and prevents bottlenecks, while efficient data transfer ensures smooth execution. The system is capable of maintaining approximately 25–30 frames per second, enabling continuous monitoring and timely alert generation.

D. System Reliability

Reliability is achieved through the use of well-established technologies across all modules. YOLO-based models provide robust object detection, MediaPipe ensures stable facial landmark extraction, and DeepFace enables consistent emotion recognition. The inclusion of YOLOv8 improves behavioral analysis through distraction detection. The modular architecture improves fault tolerance by allowing independent operation of components, reducing the impact of individual module failures. Integration of multiple sensing techniques further enhances reliability by providing complementary insights into driving conditions.

E. Expected Outcomes

The framework is expected to deliver accurate road hazard detection, reliable driver monitoring, and effective risk fusion by capturing compounded risk scenarios. It enables real-time alerts through hardware and voice assistance, reducing reaction time and improving overall road safety compared to traditional single-domain systems.

F. Advantages Over Existing Systems

Compared to conventional approaches, the framework integrates both environmental perception and driver behavior analysis into a unified architecture, enabling a comprehensive risk assessment. The risk fusion mechanism improves decision-making by capturing compounded effects rather than treating risks independently. Inclusion of drowsiness, emotion, and distraction detection improves robustness, while the real-time capability and low-cost implementation make it practical for deployment. The modular design also supports scalability and future enhancements.

G. Limitations

Despite its advantages, the framework has certain limitations. Performance may be affected under challenging conditions such as poor lighting, occlusions, or adverse weather. Detection accuracy depends on the quality and diversity of training data, particularly for the YOLOv11 model. Additionally, real-time performance may vary based on hardware capabilities, and further optimization may be required for deployment on embedded systems.

VI. CONCLUSION

This paper presented a Dual Camera Multimodal Risk Fusion System for real-time driver safety monitoring and emergency alert generation. The proposed system addresses the critical limitations of existing driver assistance solutions by simultaneously analyzing road hazards and driver behavior using a unified framework.

The system integrates a YOLOv11-based road hazard detection module with a comprehensive driver monitoring pipeline that utilizes MediaPipe for drowsiness detection, DeepFace for emotion recognition, and YOLOv8 for distraction detection. By combining these heterogeneous inputs, the system achieves a more holistic understanding of driving conditions.

A key contribution of this work is the introduction of a risk fusion engine, which models the compounded effect of multiple risk factors using a mathematical formulation. Unlike traditional approaches that consider risks independently, the proposed method amplifies the overall risk when environmental and driver-related hazards are present simultaneously, resulting in a more realistic risk assessment.

The system further incorporates a multi-layer alert mechanism, including hardware-based alerts and AI-driven voice assistance, ensuring timely and effective intervention. The use of low-cost hardware and efficient algorithms demonstrates the feasibility of deploying such systems in real-world scenarios.

In general, the proposed framework highlights the importance of multimodal integration and intelligent risk modeling in improving road safety and provides a foundation for next-generation driver assistance systems.

VII. FUTURE WORK

Although the proposed system establishes a robust multimodal framework for real-time driver safety monitoring, several technical enhancements can be pursued to improve its accuracy, scalability, and real-world applicability. Future work will focus on optimizing the YOLOv11-based road hazard detection model through the expansion of a large-scale data set, incorporating diverse environmental conditions such as varying illumination, weather dynamics, and complex traffic scenarios. Advanced training strategies, including transfer learning, data augmentation, and domain adaptation, can be used to improve generalization in heterogeneous road environments.

On the driver monitoring side, the system can be extended by integrating additional behavioral and physiological indicators such as head pose estimation, gaze tracking, facial micro-expression analysis, and steering behavior patterns, enabling a more comprehensive representation of driver cognitive state. Temporal modeling techniques, such as Long Short-Term Memory (LSTM) networks or Transformer-based architectures, can be incorporated to capture sequential dependencies in driver behavior and improve detection robustness over time.

The proposed risk fusion mechanism can be further enhanced by replacing the deterministic formulation with a learned fusion model, such as a neural network-based or probabilistic framework capable of dynamically weighting input features based on contextual relevance. In addition, incorporating uncertainty estimation techniques can improve the reliability of decisions in ambiguous scenarios.

From a system integration perspective, future work includes the implementation of Vehicle-to-Vehicle (V2V) and Vehicle-to-Infrastructure (V2I) communication, enabling distributed risk sharing and collaborative safety mechanisms. The framework can also be integrated with Advanced Driver Assistance Systems (ADAS) to enable closed-loop control actions such as adaptive braking, lane correction, and speed regulation.

Finally, optimization for deployment on embedded edge devices such as NVIDIA Jetson or similar platforms will be explored, focusing on model compression techniques that include pruning, quantization, and knowledge distillation to ensure real-time performance with minimal computational overhead. These enhancements will facilitate the transition of the proposed system from a conceptual prototype to a scalable, real-world intelligent transportation solution.

REFERENCES

- [1] S. Khare, A. Sharma, and R. Gupta, "Yolov8-based visual detection of road hazards," *arXiv preprint arXiv:2311.00073*, 2023.
- [2] A. Das, P. Roy, and S. Mukherjee, "IoT-assisted driver drowsiness detection using cnn-lstm models," *MDPI Information*, vol. 15, no. 1, pp. 1–18, 2024.

- [3] M. Espino-Salinas, J. Torres, and L. Ramirez, "Multimodal driver emotion recognition using facial expressions," *Frontiers in Artificial Intelligence*, vol. 7, 2024.
- [4] S. Venkatesan, R. Kumar, and P. Natarajan, "Human emotion detection using deepface and artificial intelligence," *MDPI Engineering Proceedings*, vol. 59, no. 1, 2023.
- [5] J. Fonseca and M. Ferreira, "Drowsiness detection: A systematic review of deep learning models," *Applied Sciences*, vol. 15, no. 16, 2025.
- [6] Q. Ji and X. Yang, "Real-time nonintrusive monitoring and prediction of driver fatigue," *IEEE Transactions on Vehicular Technology*, vol. 53, no. 4, pp. 1052–1068, 2004.
- [7] Y. Dong, Z. Hu, K. Uchimura, and N. Murayama, "Driver inattention monitoring system for intelligent vehicles," *IEEE Transactions on Intelligent Transportation Systems*, vol. 12, no. 2, pp. 596–614, 2011.
- [8] T. Soukupová and J. Čech, "Real-time eye blink detection using facial landmarks," in *21st Computer Vision Winter Workshop*, 2016.
- [9] J. Redmon, S. Divvala, R. Girshick, and A. Farhadi, "You only look once: Unified, real-time object detection," in *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 2016, pp. 779–788.
- [10] F. Zhang, X. Zhu, and M. Ye, "Mediapipe face mesh: Real-time face landmark detection and tracking," *arXiv preprint arXiv:2006.10959*, 2020.
- [11] T. Johnson, "Driver monitoring in autonomous vehicle simulation environments," *Simulation Modelling Practice and Theory*, vol. 124, 2024.
- [12] A. Nocentini, F. Rossi, and M. Conti, "Graph-based driver monitoring system using facial and skeletal features," *IEEE Access*, vol. 13, pp. 1–15, 2025.